# SMAC Manual, version 0.92b

Department of Computer Science
University of British Columbia
Vancouver, BC  V6T 1Z4, Canada.

July 24, 2012

## Contents

# 1 Introduction

SMAC (Sequential Model-based Algorithm Configuration) [CITATION NEEDED] is an SMBO [CITATION NEEDED] based automatic configuration tool for parameter optimization. It works on parameterized algorithms supporting categorical and continuous spaces with conditionality. SMAC searches through a space of possible parameter configurations by taking those which look promising, running them, updating a regression model with new information, selecting a new set of promising configurations and repeating.

Users of SMAC must provide:

- a parametric algorithm $\mathcal{A}$ (executable to be called from the command line),

- all parameters and their possible values $\Theta$, and

- a set of benchmark instances, $\Pi$.

- The objective with which to measure and aggregate algorithm preformance results.

SMAC then executes algorithm $\mathcal{A}$ with different combinations of parameters sampled from $\Theta$ with instances sampled from $\Pi$, searching for the configuration that yields overall best performance across the benchmark instances under the

supplied objective. For more details see [CITATION NEEDED]. If you use SMAC in your research, please cite that article. It would also be nice if you sent us an email – we are always interested in additional application domains.

## 1.1 History

SMAC originally grew out of improvements to Bartz-Beielstein et al's Sequential Parameter Optimization Toolbox, called SPO+ in 2009 which was fully automated and more robust than the original. In 2010 development began on a time-bounded variant which reduced computational overheads called TB-SPO. In 2011 the first version of SMAC was written in MATLAB while relying on some Java components, and in early 2012 the remainder was ported to Java.

Some of SMAC's idiosyncracies are the result of previous work on a model-free Automatic Configurator called ParamILS [CITATION NEEDED].

# 2 Configuration & Usage

This quick start guide is for **version 0.92b** of SMAC. It is important to note that this is a preliminary beta release. It has not been thoroughly tested and validated yet. Download the smac-0.90b.zip file, and unzip it in a new directory. Smac is executed by calling *smac* and by default it will merely output the usage options.

## 2.1 Requirements

SMAC itself requires nothing more than Java 1.6 to run. It has not been tested on Windows machines, and the supplied scripts are only for running under Unix, some of the examples may require other software such as CPLEX, ruby or perl. All required libraries are provided within the download package Some of the examples supplied are compiled for usage under Linux.

## 2.2 Configuration Files

Configuration Files are a way of saving different default values for frequent use with SMAC without having to specify them on every execution. Currently you cannot override values set in configuration files via the command line, though this will hopefully change. Additionally there are several configuration files that can be specified for different purposes. The general format for a configuration file is the name of the configuration option (without the two dashes) an equal sign and then the value (for booleans it should be true or false, lowercase).

### 2.2.1 Tuning Scenario Files

Perhaps the most useful of the configuration files that can be specified are *Tuning scenario files* which contain all the information about running a target algorithm. They are backwards compatible with previous versions of MATLAB

SMAC and ParamILS (to facilitate this some of the options names may be slightly different than in section (INSERT SECTION), rest assured these all are aliases of commands specified via the command line.

They can contain the following information:

**algo** An algorithm executable or a call to a wrapper script around an algorithm that conforms with the input/output format specified in section (INSERT SECTION).

**execdir** Directory to execute `<algo>` from: (e.g. "cd `<execdir>`; `<algo>`")

**deterministic** Set to 0 for randomized algorithms, 1 for deterministic .

**run_obj** Determines how to convert the resulting output line into a scalar quantifying how "good" a single algorithm execution is, (e.g. how long it took to execute, how good of a solution it found, etc...). Implemented examples for this are as follows [1]:

| Name | Description |
|------|-------------|
| **RUNTIME** | **The reported runtime of the algorithm.** |
| **RUNLENGTH** | **The reported runlength of the algorithm.** |
| **APPROX** | $1 - \dfrac{\text{optimal quality stored in instance specific information}}{\text{reported quality}}$ |
| **SPEEDUP** | $\dfrac{\text{runtime stored in instance specific information}}{\text{reported runtime}}$ |
| **QUALITY** | **The reported quality of the algorithm.** |

**overall_obj** While `run_obj` defines the objective function for a single algorithm run, `overall_obj` defines how those single objectives are combined to reach a single scalar value to compare two parameter configurations. Implemented examples for this are as follows[2]:

| Name | Description |
|------|-------------|
| **MEAN** | **The mean of the values** |
| **MEDIAN** | **The median of the values** |
| **Q90** | **The 90th percentile of the values** |
| **ADJ_MEAN** | $\dfrac{\text{total run objective}}{\# \, of \, successful \, runs}$ |
| **MEAN1000** | **Unsuccessful runs are counted as 1000 × cutoff_time** |
| **MEAN10** | **Unsuccessful runs are counted as 10 × cutoff_time** |
| **GEOMEAN** | **The geometric mean (primarily used in combination with run_obj = SPEEDUP).** |

**cutoff_time** The time after which a single algorithm execution will be terminated unsuccessfully. This is an important parameter: if choosen too high,

---

[1] APPROX, SPEEDUP rely on instance specific information which is not currently implemented

[2] ADJ_MEAN Not implemented currently

lots of time will be wasted with unsuccessful runs. If chosen too low the optimization is biased to perform well on easy instances only.

**cutoff_length** The runlength after which a single algorithm execution will be terminated unsuccesfully. The actual semantic meaning of this value is up to the target algorithm. (NOT SUPPORTED IN THIS VERSION)

**tunerTimeout** The amount of CPU time allowed for target algorithm execution (specifically once the sum of all runtimes exceeds this value, SMAC will stop configuring).

**paramfile** Specifies the file with the parameters of the algorithms. The format of this file is covered in...

**outdir** Specifies the directory ParamILS should write its results to.

**instance_file** Specifies the file with a list of training instances. Additionally this parameter can be specified as **instance_seed_file**. The format of these files is covered in...

**test_instance_file** Specifies the file with a list of test instances. Additionally this parameter can be specified as **test_instance_seed_file**. The format of these files is covered in...

### 2.2.2 Instance File Format

ParamILS and the MATLAB version of SMAC required specifying the type of file as different mutually exclusive options (e.g. **instance_file** versus **instance_seed_file**). This version detects the format automatically. These files are CSV files with " as a cell delimiter, and , as the cell seperator. If the file has 1 column, then it is an instance file, and the seeds will be randomly generated. If the file has 2 columns, then the first column must be a number and this represents a seed to be used with this instance. The second column is the filename of the instance. In the first version, each file name can be specified only once. In the second version, the instances are specified on multiple lines with duplicate entries resulting in those seeds being available for that instance. Additionally these files are also parsed with " " (space) as a cell seperator (and no cell delimiter), this is for backwards compatibility with ParamILS, note this doesn't work with spaces in filenames.

## 2.3 Configuration Options Reference

**–algoExec** Command String to execution
     Default:

**–execDir** Working directory to execute algorithm in
     Default:

**–deterministic** Whether the target algorithm is deterministic (0 no, 1 yes) [An integer due to backwards compatibility]
    Default: 0

**–targetAlgorithmEvaluator** System we should use to dispatch algorithm requests [ See manual but CLI (via Command Line) is default and generally what you want ]
    Default: CLI

**–targetAlgorithmEvaluatorSearchPath** Where we should look for other target algorithm evaluators [ See manual but generally you can ignore this ]
    Default: <current working directory> (As well as a few other folders related to internal development use)

**–runObj** Per Target Algorithm Run Objective Type that we are optimizing for
    Default: RUNTIME
    Values: RUNTIME, RUNLENGTH, APPROX, SPEEDUP, QUALITY

**–intraInstanceObj** Objective function used to aggregate multiple runs for a single instance
    Default: MEAN
    Values: MEAN, MEDIAN, Q90, ADJ_MEAN, MEAN1000, MEAN10, GEOMEAN

**–interInstanceObj** Objective function used to aggregate over multiple instances (that have already been aggregated under the Intra-Instance Objective)
    Default: MEAN
    Values: MEAN, MEDIAN, Q90, ADJ_MEAN, MEAN1000, MEAN10, GEOMEAN

**–cutoffTime** Cap Time for an Individual Run
    Default: 300.0

**–cutoffLength** Cap Time for an Individual Run [Not Implemented Currently]
    Default: -1.0

**–tunerTimeout** Total CPU Time to execute for
    Default: 2147483647

**–instanceFile** File containing instances in either "<instance filename>", or "<seed>,<instance filename>" format
    Default:

**–instanceFeatureFile** File that contains the all the instances features
    Default:

–**testInstanceFile** File containing instances specified one instance per line
    Default:

–**scenarioFile** Scenario File
    Default:

–**skipInstanceFileCheck** Do not check if instances files exist on disk
    Default: false

–**outputDirectory** Output Directory
    Default: <current working directory>/smac-output

–**paramFile** File containing Parameter Space of Execution
    Default:

–**seed** Seed for Random Number Generator (equivalent to numRun)
    Default: 0

–**experimentDir** Root Directory for Experiments Folder
    Default: <current working directory>/

–**numIterations** Total number of iterations to perform
    Default: 2147483647

–**runtimeLimit** Total Wall clock time to execute for
    Default: 2147483647

–**totalNumRunLimit** Total number of target algorithm runs to execute
    Default: 2147483647

–**runHashCodeFile** File containing a list of Run Hashes one per line (Either
    with just the format on each line, or with the following text per line: "Run
    Hash Codes: (Hash Code) After (n) runs". The number of runs in this file
    need not match the number of runs that we execute, this file only ensures
    that the sequences never diverge. Note the n is completely ignored so the
    order they are specified in is the order we expect the hash codes in this
    version
    Default:

–**modelHashCodeFile** File containing a list of Model Hashes one per line
    with the following text per line: "Preprocessed Forest Built With Hash
    Code: (n)" or "Random Forest Built with Hash Code: (n)" where (n) is
    the hashcode
    Default:

–**runGroupName** Name of subfolder of outputdir to save all outputs of this
    group (different numruns) of runs
    Default: RunGroup-2012-07-24–13-31-08-322

–**numPCA** Number of prinicipal components of features
    Default: 7

**–expectedImprovementFunction** Expected Improvement Function to Use
Default: EXPONENTIAL
Values: EXPONENTIAL, SIMPLE, SPO, EI, EIh

**–numberOfChallengers** Number of Challengers needed for Local Search
Default: 10

**–numberOfRandomConfigsInEI** Number of Random Configurations to evaluate in EI Search
Default: 10000

**–stateSerializer** Controls how the state will be saved to disk
Default: LEGACY
Values: NULL, LEGACY

**–stateDeserializer** Controls how the state will be restored from disk
Default: LEGACY
Values: NULL, LEGACY

**–restoreStateFrom** The Location (State Deserializer Dependent) of States
Default:

**–restoreIteration** The Iteration to Restore
Default:

**–executionMode** Mode of Automatic Configurator to run
Default: SMAC
Values: SMAC, ROAR

**–splitMin** Minimum number of elements needed to split a node
Default: 10

**–fullTreeBootstrap** Bootstrap all data points into trees
Default: false

**–storeDataInLeaves** Store full data in leaves of trees
Default: false

**–logModel** Store data in Log Normal form
Default: false

**–nTrees** Number of Trees in Random Forest
Default: 10

**–minVariance** Minimum allowed variance
Default: 1.0E-14

**–ratioFeatures** Number of features to consider when building Regression Forest
Default: 0.8333333333333334

8

**–preprocessMarginal** Build Random Forest with Preprocessed Marginal
Default: true

**–adaptiveCapping** Enable Adaptive Capping
Default: false

**–capSlack** Amount to scale computed cap time of challengers by
Default: 1.3

**–capAddSlack** Amount to increase computed cap time of challengers by [
general formula: capTime = capSlack*computedCapTime + capAddSlack
]
Default: 1.0

**–imputationIterations** Amount of times to impute censored data
Default: 10

**–maxConcurrentAlgoExecs** Maximum number of concurrent target algorithm executions
Default: 1

**–skipValidation** Do not perform validation at the end
Default: false

**–maxIncumbentRuns** Maximum Number of Incumbent Runs allowed
Default: 2000

**–intensificationPercentage** Percent of time to spend intensifying versus model learning
Default: 0.5

**–numSeedsPerTestInstance** Number of test seeds to use per instance during validation
Default: 1000

**–numTestInstances** Number of instances to test against (Will execute min of this, and number of instances in test Instance File)
Default: 2147483647

**–numberOfValidationRuns** Approximate Number of Validation Runs to do
Default: 1000

**–validationRoundingMode** Whether to round the number of validation runs up or down (to next multiple of numTestInstances
Default: UP
Values: UP, NONE

**–noValidationHeaders** Don't put headers on output CSV files for Validation
Default: false

**–consoleLogLevel** Default Error Level of Console Output
Default: INFO
Values: TRACE, DEBUG, INFO, WARN, ERROR

**–abortOnCrash** Treat algorithm crashes as ABORT (Useful if the algorithm really should never CRASH)
Default: false

**–abortOnFirstRunCrash** If the first run of the algorithm CRASHED treat it as an ABORT, otherwise leave it alone
Default: false

**–countSMACTimeAsTunerTime** Include the CPU Time of SMAC as part of the tunerTimeout
Default: true

**–retryTargetAlgorithmRunCount** Number of times to retry an algorithm run before reporting crashed (NOTE: The original crashes DO NOT count towards any time limits, they are in effect lost). Additionally this only retries CRASHED runs, not ABORT runs, this is by design as ABORT is only for cases when we shouldn't bother further runs
Default: 0

## 2.4 Usage

To get started with an existing configuration scenerio you simply need to execute smac as follows:

```
./smac --scenarioFile <file>
```

This will execute SMAC with the default options on the scenario specified in the file. Several of the options above can have drastic changes on the performance and usage of smac as follows:

### 2.4.1 ROAR Mode

```
./smac --scenarioFile <file> --executionMode ROAR
```

This will execute SMAC but will execute it in ROAR mode (See some paper), in short this is a model free based search where a random configuration is raced against the incumbent.

### 2.4.2 Adaptive Capping

```
./smac --scenarioFile <file> --adaptiveCapping
```

Adaptive Capping (See some paper) will cause SMAC to only schedule algorithm runs for as long as is needed to determine whether or not an algorithm is a promising configuration, as opposed to the default runtime **–cutoffTime**. This

can drastically improve tuning performance with limited time budgets. Related configuration options for this are –**capSlack**, –**capAddSlack**, –**imputationIterations**. Note: Adaptive Capping should only be used when the –**runObj** is RUNTIME.

### 2.4.3 Run Hash Codes

```
./smac --scenarioFile <file> --runHashCodeFile <logfile>
```

A Run Hash Code is a sequence of hashes that represent which runs were scheduled by SMAC. SMAC logs all Run Hash Codes to the log file, and this option allows reading of that log file for subsequent runs to ensure that the exact same set of runs is scheduled. This is primarily of use for developers.

### 2.4.4 Wall Clock Limit

```
./smac --scenarioFile <file> --runtimeLimit <seconds>
```

SMAC will abort if this amount of wall-clock time expires. This does not over ride the –**tunerTimeout** option and it's basically whichever comes first. This is useful to limit the amount of unaccounted for overhead in an algorithm run.

### 2.4.5 State Restoration

```
./smac --scenarioFile <file> --restoreStateFrom <dir> --restoreIteration <iteration>
```

SMAC will read the files in the specified directory and will change it's state such that it had previously executed all the runs specified up to the iteration specified and continue executing. Provided other options such as –**seed** and other arguments are identical you should get the exact same result out provided the target algorithm runtime variance on individual instance & seed pairs is small. This can also be used to restore runs from MATLAB (although due to the lossy nature of MATLAB files, and differences in random calls you will not get the same resulting trajectory). By default the possible states that are able to be restored are the start of iteration numbers of the form $2^n$ $n \in \mathbb{Z}_{\geq 0}$ as well as the previous 2 iterations prior to SMAC completing. If the run crashed additional information is saved, but in general you cannot restore to that iteration (since the crash most likely occured in the middle of an iteration).

    **Note:** When you restore a SMAC state, you are in essence preloading a bunch of runs and then running the scenario. In certain cases if the scenario has changed, this may result in undefined behaivor. Changing something like –**tunerTimeout** is usually a safe bet, however changing the –**runObj** may not be, as the incumbent may not actually be the best configuration under this new objective.

### 2.4.6 Checking Available Iterations To Restore

To check the available iterations that can be restored from a saved directory use:

```
./smac-possible-restores <dir>
```

### 2.4.7 Turning Off State Saving

```
./smac --scenarioFile <file> --stateSerializer NULL
```

SMAC will not save any state information to disk.

### 2.4.8 Concurrent Algorithm Execution Requests

```
./smac --scenarioFile <file> --maxConcurrentAlgoExecs <num>
```

In certain circumstances it may be much faster to allow more than one algorithm execution at once. Such as when multiple cores are available or when actual algorithm execution is I/O bound (which can happen in certain circumstances). As a result you can have SMAC schedule multiple runs at a time. In general this will result in the same trajectory as when set to 1, with the exception of when –**adaptiveCapping** is set, in this case concurrent runs are scheduled with cutoff times as if each were the first of the runs to be scheduled.

### 2.4.9 Named Runs

```
./smac --scenarioFile <file> --runID <someID>
```

All output is written to the –**outputDirectory** / –**runID**. This runID defaults to "Run-(Current Time)" a more meaningful name can be set here.

### 2.4.10 Perform Additional Validation Runs

When using instance (and not instance & seed files), SMAC Performs ~1000 (rounded up to the nearest multiple number of instances) target algorithm runs (provided there are enough instances and seeds)

```
./smac --scenarioFile <file> --runID <someID> --numValidationRuns 2000
```

### 2.4.11 Limiting the Number of Seeds Used in a Validation Run

By default SMAC limits the number of seeds used in validation runs to 1000 seeds per instance, this can be changed by using the following (This parameter has no effect on instance seed files):

```
./smac --scenarioFile <file> --runID <someID> --numValidationRuns 100000 --numSeedsPer
```

### 2.4.12 Limiting the Number of Instances Used in a Validation Run

To use only some of the instances or instance seeds specified you can limit them with the --numTestInstances parameter. For instances it will only the specified number from the top of the file, and will keep repeating them until enough seeds are used. For instance seeds it will only use the specified number of instance seeds in the file.

```
./smac --scenarioFile <file> --runID <someID> --numTestInstances 10
```

### 2.4.13 Skip Validation

```
./smac --scenarioFile <file> --runID <someID> --skipValidation
```

Does not perform any validation when after automatic configuration is completed

### 2.4.14 Manual Validating Data

SMAC also includes a method of validating data without a smac run. You can supply a configuration using the --configuration directive. Beyond this directive the rest of the options are a subset of those from SMAC.

```
./smac-validate --scenarioFile <file> --numValidationRuns 100000 --configuration "confi
```

Usage Notes:

1. If no configuration is specified the default configuration is validated.

2. This validates against the test set and not the instance set, the instance set is of no consequence here.

3. By default this outputs into the current directory, to output in the --outputDir use --useScenarioOutDir

## 3 Example tuning scenarios

The zip file contains examples for three algorithms: SAPS [?], a local search algorithm for SAT; Spear[3], a tree search algorithm for SAT (and satisfiability modulo theories) developed by Domagoj Babic; and the commercial optimization tool ILOG CPLEX [4]. Executables for Saps and Spear are included in the zip file, but CPLEX needs to be purchased to run the CPLEX example.

---

[3]http://www.cs.ubc.ca/~babic/index_spear.htm, described in some more detail in [?]
[4]http://www.ilog.com/products/cplex/

We include two tuning scenarios for Saps, one for Spear and one for CPLEX. In the Spear example, the instance collection is the same as in one of the SAPS examples: the algorithm is optimized on five graph colouring problem instances and tested on five other ones. (Note that our examples are toy examples with very small training and test data sets; we recommend substantially larger data sets for real applications!) The second tuning scenario for Saps optimizes Saps performance for a single instance, and tests on the same one, but of course with different seeds – this is useful to study peak performance of an algorithm. For CPLEX, we include a tuning scenario for mixed integer programs from combinatorial auctions (see [?] for details).

# 4  Using Instance-Specific information

Whether you choose to provide an instance_file (i.e. a list of problem instance filenames), or an instance_seed_file (i.e. a list of pairs of problem instance filenames and seeds), you specify one instance per line. You may choose to include additional information after the instance filename, such as the optimal solution quality for the instance, or the instance hardness for one or more other algorithms. Thus, the syntax for each line of the instance_file is `<instance_filename> <rest>`, where `<rest>` is an arbitrary (possibly empty) string; when using an the instance_seed_file, the syntax is `<seed> <instance_filename> <rest>` The syntax for these files allows you to do this easily: the `<rest>` string is always parsed and passed on to the objective function computation. The rest may, for example, specify a reference runtime (or runlength, or whatever) for the instance. This is very useful if the objective is to beat a competing algorithm, or a previous version of the same algorithm (In my opinion, this objective is used too much in computer science research, but since the demand is there I provide the option).

Note: `<rest>` cannot contain any whitespace currently.

# 5  Running SMAC for your own code

In order to employ SMAC to optimize your own code, you need to provide instance lists in the same format as in the above example, provide a file listing your algorithm's parameters in a predefined format, and match the required input/output format. These two latter points are covered in this section.

## 5.1  Algorithm parameter file

It is recommend you create a separate subdirectory for each algorithm you want to optimize. The parameters of your algorithm need to be defined in a file e.g. `algo-params.txt`. (Examples for such files can be found in the example directories). Comments in the file begin with a #, and run to the end of the line.

The file consists of three types of statements: parameter declarations, conditional declarations, and forbidden declarations:

**Parameter Declarations**

SMAC supports two types of parameters, continuous and categorical parameters declared as follows:

```
name { value1, ..., value_n } [defaultValue]
Example:
timeout { 1,2,4,8,16,32 } [8]
```

Here a categorical parameter is declared named timeout, it's values are one of the values listed and and the default value is 8.

```
name [minValue, maxValue] [defaultValue](i)(l)
Example 1:
timeout [1,32] [8]
Example 2:
timeout [1,32] [8]l
Example 3:
timeout [1,32] [8]i
Example 4:
timeout [1,32] [8]il
```

**Example 1** We have specified timeout as continuous with a default value of 4. Any value is legally permitted so long as it's in the interval of [1,32]. When drawing a random configuration out of this space they are drawn uniformly.

**Example 2** This is the same as Example 1, except that when drawing random configurations we do so on a log scale (e.g. a value between [1,8] is as likely to be selected as between [8,32]).

**Example 3** In this example the only legal values are integers in the range [1,32], we continue to select from the integers uniformly.

**Example 4** In this example integers in the range [1,32] are the only values permitted, and when randomly selecting them we do so on a log scale.

**Conditional Parameters**

Conditional Parameters have the following syntax:

```
dependentName | independentName operation { value1, ... , value_n}
Example Param File:
sort-algo { quick, insertion, merge, heap, stooge, bogo } [ bogo ]
quick-revert-to-insertion { 1,2,4,8,16,32 } [16]
quick-revert-to-insertion | sort-algo in { quick }
```

Currently the only supported operation is *in* , and the values are restricted to being specified as categoricals. In the above example the quick-revert-to-insertion is conditional on the sort-algo parameter being set to quick, and will be ignored otherwise.

**Forbidden Parameters**

Forbidden Parameters are parameter settings which should not be treated as valid by SMAC. During the search phase, parameters matching a forbidden parameter configuration, will not be explored. Note: Specifying a large number of these may degrade SMAC's performance.

The Syntax is as follows:

```
{name1=val1,name2=val2...}
Example Param File
quick-sort { on, off} [on]
bubble-sort { on, off} [off]
{quick-sort=on, bubble-sort=on}
{quick-sort=off, bubble-sort=off}
```

The above example implements an exclusive or (although it could be better modelled with simply one parameter), and prevents both sort techniques from being enabled at the same time, (the first forbidden parameter), and forces one to be on at all times. NOTE: The Default Parameter setting cannot itself be a forbidden parameter setting.

## 5.2  Algorithm executable / wrapper

The target algorithm as specified by the –**algoExec** parameter must obey the following general contracts. While writing your own code to achieve this is one option, there are two other common methods outlined in... (Another section)

### 5.2.1  Invocation

The algorithm must be invocable using the following:
```
<algo_executable> <instance_name> <instance_specific_information>
<cutoff_time> <cutoff_length> <seed> <param> <param> <param>...
```

**algo_executable** Exactly what is specified in the –**algoExec** argument to SMAC.

**instance_name** The name of the problem instance we are executing against

**instance_specific_information** An arbitrary string associated with this instance as specified in the –**instance_file** if no information is present then a 0 is passed here.

**cutoff_time** The amount of time in seconds that the target algorithm is permitted to run. It is the responsibility of the target algorithm to ensure that this is obeyed. It is not necessary that that the actual algorithm execution time (wall clock time) be below this value. If for instance clean up is needed, or it is only possible for the algorithm to abort at certain stages.

**cutoff_length** A domain (target algorithm) specific measure of when the algorithm should consider itself done.

**seed** A positive integer that the algorithm should use to seed itself (for reproducability), -1 is used when the algorithm is –**deterministic.**

**param** A setting of an active parameter for the selected configuration as specified in the Algorithm Parameter File. SMAC will only pass parameters that are considered active, and SMAC is not guaranteed to pass the parameters in any particular order. The exact format for each parameter is:

```
-name 'value'
```

All the parameters above are mandatory, even if they are inapplicable, in which case a dummy value will be passed which you are free to ignore.

### 5.2.2 Output

The Target Algorithm is free to output anything, which will be ignored but must at some point output a line (only once) in the following format[5]:

```
    Result for ParamILS: <solved>, <runtime>, <runlength>, <quality>,
<seed>
```

**solved** Must be one of **SAT** (signifying a successful run with a result of YES), **UNSAT** (signifying a successful run with a result of NO), **TIMEOUT** if the algorithm didn't complete, **CRASHED** if something untoward happened during the algorithm run, or **ABORT** if something prevents the target algorithm for successfully executing and it is believed that further attempts would be futile. SMAC does not differentiate between **SAT** and **UNSAT** responses, and the primary use of these is to serve as a check that the algorithm is executing correctly by allowing the algorithm to output a yes / no output that can be checked. This can be useful for debugging purposes. **ABORT** can be useful in cases where the target algorithm cannot find required files, or a permission problem prevents access to them. This will case SMAC to stop running immediately.

---

[5]ParamILS in not a typo. While other values are possible including SMAC, HAL. ParamILS is probably the most portable. The Exactly Regex that is used in this version is: **(Final)?\s\*[Rr]esult\s+(?:(for)|(of))\s+(?:(HAL)|(ParamILS)|(SMAC)|(this wrapper))**

**runtime** The amount of CPU time used during this algorithm run. SMAC does not measure the CPU time directly, and this is the amount that is used with respect to **–tunerTimeout**. You may get unexpected performance degredation when this amount is heavily under reported (this typically happens when targetting very short algorithm runs with large overheads that aren't accounted for).
**NOTE:** The **runtime** should always be strictly less than the requested **cutoff_time** when reporting **SAT** or **UNSAT**. It should always be the **cutoff_time** when reporting **TIMEOUT**. This is even the case if the algorithm can determine that it will **TIMEOUT** without doing the actual work.

**runlength** A domain specific measure of how far the algorithm progressed.

**quality** A domain specific measure of the quality of the solution.

**seed** The seed value that was used in this target algorithm execution. Note: This seed **MUST** match the seed that the algorithm was called with. This is used as fail-safe check to ensure that the output we are parsing really matches the call we requested.

Like invocation, all fields are mandatory, when not applicable 0's can be substituted.

### 5.2.3 Wrappers & Native Libraries

Beyond simplying adding the additional functionality to your code by hand, there exist several other approaches from related projects.

**Wrappers** Executable Scripts that manage the resource limits automatically and format the specified string into something usable by the actual target algorithm. This approach is probably the most common, but among it's draw backs are the fact that they often rely on third party languages, and for smaller execution times have a large amount of overhead that may not be accounted for. Most of the examples included in SMAC use this approach, and the wrappers included can be adapted for your own project

**Native Libraries** Libraries exist (See: ...) for C and Java currently that facilitate adding the functionality directly to the code, while parsing the arguments into the necessary data structures is still required, they do ensure that the output is written properly in most cases. Draw backs include not necessarily outputting the values in certain crashes.

# 6   Interpreting SMAC's Output

SMAC outputs a variety of information to log files, trajectory files, and state files. Most of the files are human readable, and this section governs where to find information.

## 6.1 Logs

SMAC's output, especially in this release is particularly verbose, and perhaps noisey. By default SMAC uses slf4j for logging, and specifically ships with logback. In theory this allows a much richer configuration option by modifying `conf/logback.xml` . To change the output level (which supports TRACE, DEBUG, INFO, WARN, ERROR, edit this file. Future versions will include more simple flexibilty, but the default configuration outputs three files into the **–outputDirectory**.

**log.txt** A formatted listing of the output of SMAC with information useful for debugging and tracing.

**log-warn.txt** Contains the same information as the above file, except only form warning and higher levels

**raw.txt** The raw messages logged by the program, this is useful for grepping.

**runhashes.txt** A file that contains the Run Hash Codes generated by this run

By default the standard output of the SMAC is in between log.txt and raw.txt. For the most part the same lines are logged however. *Some* of the output from the log file was written to be identical to the MATLAB version, but by and large this should not be relied upon as neither the MATLAB version nor this version have any well defined guarantees on what is or isn't outputted, or how. Any mechanisms relying on this are more than likely very brittle, especially as changes to the logback.xml file may result in changes to the formatting of this file.

### 6.1.1 Interpretting the Log File

SMAC execution basically consists of two stages, setup and then the actual algorithm. In the setup process alot of validation is done to hopefully catch alot of errors. Any warnings that are outputted in the first part of SMAC's run may deserve special attention, and can in some cases be turned off if they are potentially false positives. SMAC also outputs it's configuration for the run at the start of every log file. Eventually SMAC will start running, and this occurs when you see something along the lines of:

```
Automatic Configuration Start Time is May 23, 2012
```

At this point SMAC has started the actual algorithm configuration as outlined in (Insert REF). The two most important pieces of information for the general user, are probably lines like the following:

```
At end of iteration 2, incumbent is 25433[4.0, 2.0, 1.0, 7.0] -alpha '1.189' -ps '0.033
*****Runtime Statistics*****
Iteration: 2
```

19

```
Wallclock Time: 8.317 s
Wallclock Time Remaining:2.147483638683E9 s
Total CPU Time: 7.679999999999996 s
CPU Time Remaining: 22.320000000000004 s
AC CPU Time: 0.75 s
AC User Time: 0.72 s
Max Memory: 5365.375 MB
Total Java Memory: 361.5 MB
Free Java Memory: 318.36956787109375 MB
```

The first line (after the closing ]) advises you what the incumbent was at this point of time. The next lines give you the state of SMAC. Most numbers are straight forward except the AC CPU Time: and AC User Time: These are the CPU and User times of SMAC itself. CPU Time remaining is the the −**tunerTimeout** - The Total CPU Time So far.

In version **0.92b** and later, another file log-warn.txt was created. This contains only warning and error messages and is probably easier to read than the standard log.txt file. Currently some warnings will be displayed as follows:

```
Target Algorithm Runs have a hard coded quality and run length
```

This can safely be ignored, and is a reminder that this version does not support these features. Additionally if you reuse the same runID, you may get warnings that previous state output was found and renamed.

### 6.1.2 State Files

The state files stored in the state/ subdirectory of the −**outputDirectory** are a bit more well defined. The txt and csv files are human readable. The **paramstrings-itN.txt** file correspond to a listing of parameter strings. The number in the front corresponds to the numbers used in the **runs_and_results-itN.csv** file

The columns of this file in order represent:

1. The number of the row (run)

2. The number of the configuration in the paramstrings-itN.txt file

3. The number of the instance as specified in −**instanceFile**

4. The response value (−**runObjective**) of the run

5. 1 if the run was censored, 0 if not.

6. The cutoff time of the run

7. The resulting seed of the run.

8. The runtime of the run.

9. The runlength of the run

10. 1 if the run was a successful run (SAT or UNSAT), 0 otherwise.

11. The quality of the run.

12. The iteration the run occured in.

13. The cumulative sum of all the runs up to this point

### 6.1.3 Trajectory File

SMAC also outputs a trajectory file, which is a CSV file that outpus the following columns

| Column Name | Description |
|---|---|
| Total Time | Sum of all execution times and CPU time of SM |
| Incumbents Mean Performance | Performance of the Incumbent under the given −**runObjective** and |
| Incumbent's Performance $\sigma$ | Standard deviation of the Incumbent under the given −**runObjective** |
| Incumbent ID | The ID of the incumbent as listed in the paramstrings file whe |
| acTime | Total CPU Time of SMAC |
| Param String Columns | The rest give the parameter strings of the incumb |

### 6.1.4 Validation Output

When Validating Three output files are generated:

1. RawValidationExecutionResults.csv - CSV File containing a list of seeds
   & instances and the result of the wrapper execution (useful for debugging)

2. validationResultsList.csv - CSV File containing a list of seeds & instances
   and the resulting response

3. validationResultsMatrix.csv - CSV File containing the list of instances on
   each line, the second column is the aggregation of the remaining columns
   which is the response value of the run, aggergated under the −**overallObjective**.
   Finally there is one additional row that gives the aggregation of all the
   individual −**overallObjectives**, aggregated in the same way.

# 7 Usage Differences Between This and Other Versions

This section is primarily for users of ParamILS, but also the MATLAB version
of SMAC and outlines some of the differences.

## 7.1 Configuration Changes

Everything needed to execute SMAC is specified on the Command Line. Certain configuration options allow you to specify files, for instance a scenario file. While most existing scenario files should work without a hitch it's important to note that what is actually happening is that it is just being parsed into Command Line options. You could in theory specify other configuration options in a scenario file, and they would be set accordingly.

## 7.2 Previous Features Not Supported

Extended Trajectory Files from MATLAB SMAC do not exist.

# 8 Version Information

## 8.1 License

SMAC is to be released under a dual usage license. Academic & Non-commercial usage is permitted. Commercial Usage requires a license.

## 8.2 Change Log

Version History
  −Version 0.92b −
  Third Beta Release: June 16th 2012

- Added support for ABORT response from wrapper

- Fixed spurious error condition about conflicting instances being loaded when test instances and training instances aren't disjoint

- Implemented a new log file that displays only warnings or errors

- Added support for Run Hash Code File

- Fixed an error message that did not display affected filenames when it should have

- Added additional internal input validation for robustness

- Added additional compatibility fixes for restoring state from MATLAB

- Many fixes for Deterministic Algorithm Runs, including saving and restoring state, and seed selection

- Fixed issue with Local Search choosing poor paths and potentially getting into an infinite loop

- Limited the number of runs we attempt on the incumbent to the number of instance seeds available

- Forced all instances to have the same number of seeds specified (only for instance file, not instance seed file)

–Version 0.91b –
Second Beta Release: June 8th 2012

- Fixed NullPointerException when crashing on iteration 0

- Added Support for Instance Specific Information

- Added Validation Support and Utility

- Provided Domain of Configuration Options for SMAC

- Validated Paremeter Setting for Default Configuration

- Cleaned Up Error Messages related to Parsing Param File

- Added Support for Forbidden Parameters

– Version 0.90b –
Initial (Internal) Release: May 25th 2012

## 8.3   Known problems

1. You cannot override options on the command line specified in files

2. Validation of configuration options is poor and limited

3. Neighbourhood search of discretized continuous parameters may select identical parameter configurations.

4. Cutoff Length is not supported

5. Standard Deviation not reported in Trajectory Files

6. **–tunerTimeout** is computed simply on target algorithm runtimes, it may not agree with ParamILS nor MATLAB SMAC's definition

7. Scripts do not work outside of smac directory.

# References

document